



# MedBLIP – Image Captioning for Radiology X- Ray Images

depi final project  
name: dalia refaat

# MedBLIP - Image Captioning for Radiology X-Ray Images

## Overview:

This repository implements BLIP (Bootstrapping Language-Image Pre-training) to generate captions from Radiology X-ray images.

## What is BLIP?

A powerful vision-language pre-training model designed for efficient handling of image-captioning tasks.

## Focus of the Project:

The project specifically targets medical imaging, applying BLIP to the Radiology Objects in Context (ROCO) dataset.

## Importance:

Enhances the accessibility of medical imaging by generating descriptive captions, aiding in better understanding and interpretation of X-ray images.

The BLIP (Bootstrapping Language-Image Pre-training) model comprises several components designed to enhance its performance in vision-language tasks. Here's a breakdown of the models and components within BLIP:

1. Vision Encoder:

Architecture: Typically based on Convolutional Neural Networks (CNNs) or Vision Transformers (ViTs).

Function: This component extracts features from the input images, providing a rich representation that the language model can utilize.

2. Language Decoder:

Architecture: Often implemented using Transformer architectures, similar to those used in models like GPT or BERT.

Function: This decoder takes the encoded visual features and generates text captions based on those features.

3. Bootstrapping Mechanism:

Function: This unique aspect of BLIP allows it to iteratively improve its language-image representations. It leverages both image and text data during training to enhance the understanding of context and meaning.

4. Cross-Modal Attention:

Function: The attention layers enable the model to focus on specific regions of an image while generating corresponding text. This is crucial for generating accurate and contextually relevant captions.

5. Multi-Task Learning:

Function: BLIP can be trained on various tasks, such as image captioning, visual question answering, and image-text retrieval, which enhances its versatility.

6. Pre-trained Checkpoints:

Function: BLIP often provides pre-trained models on large datasets, allowing users to fine-tune them on specific tasks or domains, such as medical imaging, without needing to train from scratch.

Summary of Models:

Vision Encoder: Extracts visual features (CNNs or ViTs).

Language Decoder: Generates textual captions (Transformers).

Bootstrapping Mechanism: Enhances representations iteratively.

Cross-Modal Attention: Links visual features with text.

Multi-Task Learning: Trained for various vision-language tasks.

Pre-trained Checkpoints: Ready for fine-tuning on specialized datasets.

# About dataset:

## ROCO-Dataset

Data Card Code (18) Discussion (1) Suggestions (0)

### About Dataset

No description available

**all\_data** (3 directories)

⬇ ⌵ ⬆

#### About this directory

📝 Add Suggestion

This file does not have a description yet.

test  
2 directories, 1  
files

train  
2 directories, 1  
files

validation  
2 directories, 1  
files

#### Usability ⓘ

2.94

#### License

CC0: Public Domain

#### Expected update frequency

Not specified

#### Tags

#### Data Explorer

Version 1 (7.77 GB)

- all\_data
  - test
  - train
  - validation

#### Summary

- 87.9k files
- 18 columns



## Data Explorer

Version 1 (7.77 GB)

- all\_data
  - test
    - non-radiology
    - radiology
      - radiologytestdata.csv
  - train
    - non-radiology
    - radiology
      - images
        - captions.txt
        - cuis.txt
        - dlinks.txt
        - keywords.txt
        - licences.txt
        - semtypes.txt
      - traindata.csv
      - radiologytraindata.csv
    - validation
      - non-radiology
      - radiology
        - radiologyvaldata.csv

## Summary

- 87.9k files
- 18 columns


# checking for missing values:

Explanation:

- This function checks for missing or invalid image files in the DataFrame.
- It iterates through the image paths, verifying if each file exists and can be opened.
- If invalid files are found, those rows are removed from the DataFrame to ensure data integrity.

## 3. Checking for Missing or Invalid Files

python

 Copy code

```
def check_and_remove_missing_or_invalid_files(df, column_name):  
    missing_or_invalid_files = []  
  
    for idx, image_path in df[column_name].items():  
        if not os.path.isfile(image_path): # Check if file exists  
            missing_or_invalid_files.append(idx)  
        else:  
            try:  
                img = Image.open(image_path)  
                img.verify() # Verify that it is, indeed, an image  
            except Exception as e:  
                print(f"Error with file {image_path}: {e}")  
                missing_or_invalid_files.append(idx)  
  
    if missing_or_invalid_files:  
        print(f"Removing {len(missing_or_invalid_files)} rows with missing or in  
        df = df.drop(missing_or_invalid_files).reset_index(drop=True)  
    else:  
        print("All files are present and valid!")  
  
    return df
```

# Dataset Class creation:

Explanation:

- The ImageCaptioningDataset class inherits from Dataset, allowing for a customized dataset structure.
- The \_\_init\_\_ method initializes the dataset, processor, and image size.
- The \_\_len\_\_ method returns the total number of samples in the dataset.
- The \_\_getitem\_\_ method processes the image and caption, preparing them for input to the model.

## 4. Dataset Class Creation

python

Copy code

```
class ImageCaptioningDataset(Dataset): # Custom Dataset class for image caption
    def __init__(self, dataset, processor, image_size=(224, 224)):
        self.dataset = dataset
        self.processor = processor
        self.image_size = image_size
        self.resize_transform = Resize(image_size)

    def __len__(self):
        return len(self.dataset) # Returns the number of samples in the dataset

    def __getitem__(self, idx):
        item = self.dataset.iloc[idx] # Get the sample at index idx
        img = Image.open(item['images']) # Open the image file
        encoding = self.processor(images=img, text=item['caption'], padding="max")
        # Process the image and text

        encoding = {k: v.squeeze() for k, v in encoding.items()} # Adjust encoding
        return encoding # Return processed data
```



# Checkpoint Management Functions:

## Explanation:

- The load\_checkpoint function checks if a checkpoint file exists.
- If it does, the model and optimizer states are restored, allowing training to resume from where it left off.
- If not, it initializes the training from the beginning, setting the starting epoch to zero.
- The save\_checkpoint function creates a directory for saving checkpoints if it doesn't exist.
- It saves the current epoch, model state, optimizer state, and average loss to a file.
- This function ensures that progress can be restored later without losing any training data.

## 11. Checkpoint Management Functions

### Load Checkpoint Function

python

 Copy code

```
def load_checkpoint(model, optimizer, checkpoint_path):
    if os.path.exists(checkpoint_path):
        checkpoint = torch.load(checkpoint_path) # Load the checkpoint file
        model.load_state_dict(checkpoint['model_state_dict']) # Load the model
        optimizer.load_state_dict(checkpoint['optimizer_state_dict']) # Load the optimizer
        start_epoch = checkpoint['epoch'] # Get the starting epoch from the checkpoint
        print(f"Loaded checkpoint from {checkpoint_path} at epoch {start_epoch}")
    else:
        start_epoch = 0 # If no checkpoint, start from epoch 0
        print(f"No checkpoint found at {checkpoint_path}. Starting from scratch.")
    return start_epoch
```

### Save Checkpoint Function

python

 Copy code

```
def save_checkpoint(model, optimizer, epoch, avg_loss, checkpoint_path='checkpoints/'):
    if not os.path.exists('checkpoints'):
        os.makedirs('checkpoints') # Create checkpoints directory if it doesn't exist
    torch.save({
        'epoch': epoch,
        'model_state_dict': model.state_dict(), # Save the model state
        'optimizer_state_dict': optimizer.state_dict(), # Save the optimizer state
        'avg_loss': avg_loss, # Save the average loss
    }, checkpoint_path) # Save checkpoint file
    print(f"Checkpoint saved at epoch {epoch} with average loss: {avg_loss:.4f}")
```

# Loss Calculation and Backpropagation:

Explanation:

- The loss calculated from the model's output is used for backpropagation.
- The optimizer clears previous gradients, computes the new gradients based on the current loss, and updates the model parameters accordingly.

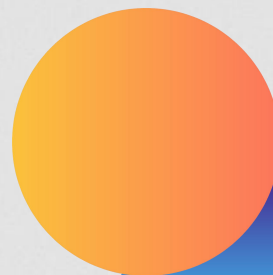
## 9. Loss Calculation and Backpropagation

python

 Copy code

```
loss = outputs.loss # Get loss

# Backpropagation
optimizer.zero_grad() # Clear previous gradients
loss.backward() # Backpropagate the loss
optimizer.step() # Update model parameters
```






# BLEU Score Calculation Function:

Explanation:

- The `calculate_bleu_score` function computes the BLEU score, a metric for evaluating the quality of generated text.
- It splits both reference captions and generated captions into words before calculating the score.
- This function can be used to assess the quality of captions generated by the model.

## 13. BLEU Score Calculation Function

python

 Copy code

```
def calculate_bleu_score(references, hypothesis):  
    references = [ref.split() for ref in references] # Split reference sentence  
    hypothesis = hypothesis.split() # Split hypothesis sentence into words  
    score = sentence_bleu(references, hypothesis) # Calculate BLEU score  
    return score # Return the BLEU score
```

